

225229108

DINESH KUMAR

Importing csv file**STEP1:**In [1]: `import pandas as pd`

In [2]: `data=pd.read_csv('diabetes.csv')`
`data`

Out[2]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|--|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | 0.627 | |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | 0.351 | |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | |
| 5 | 5 | 116 | 74 | 0 | 0 | 25.6 | 0.201 | |
| 6 | 3 | 78 | 50 | 32 | 88 | 31.0 | 0.248 | |
| 7 | 10 | 115 | 0 | 0 | 0 | 35.3 | 0.134 | |
| 8 | 2 | 197 | 70 | 45 | 543 | 30.5 | 0.158 | |
| 9 | 8 | 125 | 96 | 0 | 0 | 0.0 | 0.232 | |
| 10 | 4 | 110 | 92 | 0 | 0 | 37.6 | 0.191 | |
| 11 | 10 | 168 | 74 | 0 | 0 | 38.0 | 0.537 | |
| 12 | 10 | 139 | 80 | 0 | 0 | 27.1 | 1.441 | |
| 13 | 1 | 189 | 60 | 23 | 846 | 30.1 | 0.398 | |
| 14 | 5 | 166 | 72 | 19 | 175 | 25.8 | 0.587 | |
| 15 | 7 | 100 | 0 | 0 | 0 | 30.0 | 0.484 | |
| 16 | 0 | 118 | 84 | 47 | 230 | 45.8 | 0.551 | |
| 17 | 7 | 107 | 74 | 0 | 0 | 29.6 | 0.254 | |
| 18 | 1 | 103 | 30 | 38 | 83 | 43.3 | 0.183 | |
| 19 | 1 | 115 | 70 | 30 | 96 | 34.6 | 0.529 | |
| 20 | 3 | 126 | 88 | 41 | 235 | 39.3 | 0.704 | |
| 21 | 8 | 99 | 84 | 0 | 0 | 35.4 | 0.388 | |
| 22 | 7 | 196 | 90 | 0 | 0 | 39.8 | 0.451 | |
| 23 | 9 | 119 | 80 | 35 | 0 | 29.0 | 0.263 | |
| 24 | 11 | 143 | 94 | 33 | 146 | 36.6 | 0.254 | |
| 25 | 10 | 125 | 70 | 26 | 115 | 31.1 | 0.205 | |
| 26 | 7 | 147 | 76 | 0 | 0 | 39.4 | 0.257 | |
| 27 | 1 | 97 | 66 | 15 | 140 | 23.2 | 0.487 | |
| 28 | 13 | 145 | 82 | 19 | 110 | 22.2 | 0.245 | |
| 29 | 5 | 117 | 92 | 0 | 0 | 34.1 | 0.337 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 738 | 2 | 99 | 60 | 17 | 160 | 36.6 | 0.453 | |
| 739 | 1 | 102 | 74 | 0 | 0 | 39.5 | 0.293 | |

| Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | Diabetes | Pedigree | Function |
|-------------|---------|---------------|---------------|---------|-----|----------|----------|----------|
| 740 | 11 | 120 | 80 | 37 | 150 | 42.3 | | 0.785 |
| 741 | 3 | 102 | 44 | 20 | 94 | 30.8 | | 0.400 |
| 742 | 1 | 109 | 58 | 18 | 116 | 28.5 | | 0.219 |
| 743 | 9 | 140 | 94 | 0 | 0 | 32.7 | | 0.734 |
| 744 | 13 | 153 | 88 | 37 | 140 | 40.6 | | 1.174 |
| 745 | 12 | 100 | 84 | 33 | 105 | 30.0 | | 0.488 |
| 746 | 1 | 147 | 94 | 41 | 0 | 49.3 | | 0.358 |
| 747 | 1 | 81 | 74 | 41 | 57 | 46.3 | | 1.096 |
| 748 | 3 | 187 | 70 | 22 | 200 | 36.4 | | 0.408 |
| 749 | 6 | 162 | 62 | 0 | 0 | 24.3 | | 0.178 |
| 750 | 4 | 136 | 70 | 0 | 0 | 31.2 | | 1.182 |
| 751 | 1 | 121 | 78 | 39 | 74 | 39.0 | | 0.261 |
| 752 | 3 | 108 | 62 | 24 | 0 | 26.0 | | 0.223 |
| 753 | 0 | 181 | 88 | 44 | 510 | 43.3 | | 0.222 |
| 754 | 8 | 154 | 78 | 32 | 0 | 32.4 | | 0.443 |
| 755 | 1 | 128 | 88 | 39 | 110 | 36.5 | | 1.057 |
| 756 | 7 | 137 | 90 | 41 | 0 | 32.0 | | 0.391 |
| 757 | 0 | 123 | 72 | 0 | 0 | 36.3 | | 0.258 |
| 758 | 1 | 106 | 76 | 0 | 0 | 37.5 | | 0.197 |
| 759 | 6 | 190 | 92 | 0 | 0 | 35.5 | | 0.278 |
| 760 | 2 | 88 | 58 | 26 | 16 | 28.4 | | 0.766 |
| 761 | 9 | 170 | 74 | 31 | 0 | 44.0 | | 0.403 |
| 762 | 9 | 89 | 62 | 0 | 0 | 22.5 | | 0.142 |
| 763 | 10 | 101 | 76 | 48 | 180 | 32.9 | | 0.171 |
| 764 | 2 | 122 | 70 | 27 | 0 | 36.8 | | 0.340 |
| 765 | 5 | 121 | 72 | 23 | 112 | 26.2 | | 0.245 |
| 766 | 1 | 126 | 60 | 0 | 0 | 30.1 | | 0.349 |
| 767 | 1 | 93 | 70 | 31 | 0 | 30.4 | | 0.315 |

768 rows × 9 columns



In [3]: `data.head()`

Out[3]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-------|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | | 0.627 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | | 0.351 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | | 0.672 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | | 0.167 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | | 2.288 |

In [5]: `data.shape`

Out[5]: (768, 9)

In [6]: `data.columns`

Out[6]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'], dtype='object')

In [7]: `data.dtypes`

Out[7]: Pregnancies int64
Glucose int64
BloodPressure int64
SkinThickness int64
Insulin int64
BMI float64
DiabetesPedigreeFunction float64
Age int64
Outcome int64
dtype: object

In [8]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
Pregnancies          768 non-null int64
Glucose              768 non-null int64
BloodPressure        768 non-null int64
SkinThickness        768 non-null int64
Insulin              768 non-null int64
BMI                  768 non-null float64
DiabetesPedigreeFunction 768 non-null float64
Age                  768 non-null int64
Outcome              768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [12]: `data.BloodPressure.value_counts`

Out[12]: <bound method IndexOpsMixin.value_counts of 0>

| | |
|-----|----|
| 1 | 66 |
| 2 | 64 |
| 3 | 66 |
| 4 | 40 |
| 5 | 74 |
| 6 | 50 |
| 7 | 0 |
| 8 | 70 |
| 9 | 96 |
| 10 | 92 |
| 11 | 74 |
| 12 | 80 |
| 13 | 60 |
| 14 | 72 |
| 15 | 0 |
| 16 | 84 |
| 17 | 74 |
| 18 | 30 |
| 19 | 70 |
| 20 | 88 |
| 21 | 84 |
| 22 | 90 |
| 23 | 80 |
| 24 | 94 |
| 25 | 70 |
| 26 | 76 |
| 27 | 66 |
| 28 | 82 |
| 29 | 92 |
| | .. |
| 738 | 60 |
| 739 | 74 |
| 740 | 80 |
| 741 | 44 |
| 742 | 58 |
| 743 | 94 |
| 744 | 88 |
| 745 | 84 |
| 746 | 94 |
| 747 | 74 |
| 748 | 70 |
| 749 | 62 |
| 750 | 70 |
| 751 | 78 |
| 752 | 62 |
| 753 | 88 |
| 754 | 78 |
| 755 | 88 |
| 756 | 90 |
| 757 | 72 |
| 758 | 76 |
| 759 | 92 |
| 760 | 58 |
| 761 | 74 |

```
762    62
763    76
764    70
765    72
766    60
767    70
Name: BloodPressure, Length: 768, dtype: int64>
```

Step-2.[Build Logistic Regression model]

```
In [13]: X=data[['Pregnancies','Glucose','BloodPressure','SkinThickness','Insulin','BMI','|  
In [15]: Y=data[['Outcome']]  
In [16]: from sklearn.model_selection import train_test_split  
In [18]: X_train,X_test,y_train,y_test=train_test_split(X,Y,test_size=0.25,random_state=42|
```

In [19]: X_train

Out[19]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|--|
| 357 | 13 | 129 | 0 | 30 | 0 | 39.9 | 0.569 | |
| 73 | 4 | 129 | 86 | 20 | 270 | 35.1 | 0.231 | |
| 352 | 3 | 61 | 82 | 28 | 0 | 34.4 | 0.243 | |
| 497 | 2 | 81 | 72 | 15 | 76 | 30.1 | 0.547 | |
| 145 | 0 | 102 | 75 | 23 | 0 | 0.0 | 0.572 | |
| 514 | 3 | 99 | 54 | 19 | 86 | 25.6 | 0.154 | |
| 291 | 0 | 107 | 62 | 30 | 74 | 36.6 | 0.757 | |
| 132 | 3 | 170 | 64 | 37 | 225 | 34.5 | 0.356 | |
| 559 | 11 | 85 | 74 | 0 | 0 | 30.1 | 0.300 | |
| 631 | 0 | 102 | 78 | 40 | 90 | 34.5 | 0.238 | |
| 719 | 5 | 97 | 76 | 27 | 0 | 35.6 | 0.378 | |
| 395 | 2 | 127 | 58 | 24 | 275 | 27.7 | 1.600 | |
| 41 | 7 | 133 | 84 | 0 | 0 | 40.2 | 0.696 | |
| 637 | 2 | 94 | 76 | 18 | 66 | 31.6 | 0.649 | |
| 108 | 3 | 83 | 58 | 31 | 18 | 34.3 | 0.336 | |
| 481 | 0 | 123 | 88 | 37 | 0 | 35.2 | 0.197 | |
| 56 | 7 | 187 | 68 | 39 | 304 | 37.7 | 0.254 | |
| 323 | 13 | 152 | 90 | 33 | 29 | 26.8 | 0.731 | |
| 685 | 2 | 129 | 74 | 26 | 205 | 33.2 | 0.591 | |
| 758 | 1 | 106 | 76 | 0 | 0 | 37.5 | 0.197 | |
| 572 | 3 | 111 | 58 | 31 | 44 | 29.5 | 0.430 | |
| 529 | 0 | 111 | 65 | 0 | 0 | 24.6 | 0.660 | |
| 24 | 11 | 143 | 94 | 33 | 146 | 36.6 | 0.254 | |
| 465 | 0 | 124 | 56 | 13 | 105 | 21.8 | 0.452 | |
| 247 | 0 | 165 | 90 | 33 | 680 | 52.3 | 0.427 | |
| 443 | 8 | 108 | 70 | 0 | 0 | 30.5 | 0.955 | |
| 351 | 4 | 137 | 84 | 0 | 0 | 31.2 | 0.252 | |
| 327 | 10 | 179 | 70 | 0 | 0 | 35.1 | 0.200 | |
| 110 | 3 | 171 | 72 | 33 | 135 | 33.3 | 0.199 | |
| 82 | 7 | 83 | 78 | 26 | 71 | 29.3 | 0.767 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 21 | 8 | 99 | 84 | 0 | 0 | 35.4 | 0.388 | |
| 313 | 3 | 113 | 50 | 10 | 85 | 29.5 | 0.626 | |
| 459 | 9 | 134 | 74 | 33 | 60 | 25.9 | 0.460 | |

| Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | Diabetes | PedigreeFunction |
|-------------|---------|---------------|---------------|---------|-----|----------|------------------|
| 160 | 4 | 151 | 90 | 38 | 0 | 29.7 | 0.294 |
| 276 | 7 | 106 | 60 | 24 | 0 | 26.5 | 0.296 |
| 191 | 9 | 123 | 70 | 44 | 94 | 33.1 | 0.374 |
| 385 | 1 | 119 | 54 | 13 | 50 | 22.3 | 0.205 |
| 413 | 1 | 143 | 74 | 22 | 61 | 26.2 | 0.256 |
| 491 | 2 | 89 | 90 | 30 | 0 | 33.5 | 0.292 |
| 343 | 5 | 122 | 86 | 0 | 0 | 34.7 | 0.290 |
| 308 | 0 | 128 | 68 | 19 | 180 | 30.5 | 1.391 |
| 661 | 1 | 199 | 76 | 43 | 0 | 42.9 | 1.394 |
| 130 | 4 | 173 | 70 | 14 | 168 | 29.7 | 0.361 |
| 663 | 9 | 145 | 80 | 46 | 130 | 37.9 | 0.637 |
| 99 | 1 | 122 | 90 | 51 | 220 | 49.7 | 0.325 |
| 372 | 0 | 84 | 64 | 22 | 66 | 35.8 | 0.545 |
| 87 | 2 | 100 | 68 | 25 | 71 | 38.5 | 0.324 |
| 458 | 10 | 148 | 84 | 48 | 237 | 37.6 | 1.001 |
| 330 | 8 | 118 | 72 | 19 | 0 | 23.1 | 1.476 |
| 214 | 9 | 112 | 82 | 32 | 175 | 34.2 | 0.260 |
| 466 | 0 | 74 | 52 | 10 | 36 | 27.8 | 0.269 |
| 121 | 6 | 111 | 64 | 39 | 0 | 34.2 | 0.260 |
| 614 | 11 | 138 | 74 | 26 | 144 | 36.1 | 0.557 |
| 20 | 3 | 126 | 88 | 41 | 235 | 39.3 | 0.704 |
| 700 | 2 | 122 | 76 | 27 | 200 | 35.9 | 0.483 |
| 71 | 5 | 139 | 64 | 35 | 140 | 28.6 | 0.411 |
| 106 | 1 | 96 | 122 | 0 | 0 | 22.4 | 0.207 |
| 270 | 10 | 101 | 86 | 37 | 0 | 45.6 | 1.136 |
| 435 | 0 | 141 | 0 | 0 | 0 | 42.4 | 0.205 |
| 102 | 0 | 125 | 96 | 0 | 0 | 22.5 | 0.262 |

576 rows × 8 columns



In [20]: X_test

Out[20]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|
| 668 | 6 | 98 | 58 | 33 | 190 | 34.0 | 0.430 |
| 324 | 2 | 112 | 75 | 32 | 0 | 35.7 | 0.148 |
| 624 | 2 | 108 | 64 | 0 | 0 | 30.8 | 0.158 |
| 690 | 8 | 107 | 80 | 0 | 0 | 24.6 | 0.856 |
| 473 | 7 | 136 | 90 | 0 | 0 | 29.9 | 0.210 |
| 204 | 6 | 103 | 72 | 32 | 190 | 37.7 | 0.324 |
| 97 | 1 | 71 | 48 | 18 | 76 | 20.4 | 0.323 |
| 336 | 0 | 117 | 0 | 0 | 0 | 33.8 | 0.932 |
| 568 | 4 | 154 | 72 | 29 | 126 | 31.3 | 0.338 |
| 148 | 5 | 147 | 78 | 0 | 0 | 33.7 | 0.218 |
| 667 | 10 | 111 | 70 | 27 | 0 | 27.5 | 0.141 |
| 212 | 7 | 179 | 95 | 31 | 0 | 34.2 | 0.164 |
| 199 | 4 | 148 | 60 | 27 | 318 | 30.9 | 0.150 |
| 265 | 5 | 96 | 74 | 18 | 67 | 33.6 | 0.997 |
| 760 | 2 | 88 | 58 | 26 | 16 | 28.4 | 0.766 |
| 356 | 1 | 125 | 50 | 40 | 167 | 33.3 | 0.962 |
| 501 | 3 | 84 | 72 | 32 | 0 | 37.2 | 0.267 |
| 457 | 5 | 86 | 68 | 28 | 71 | 30.2 | 0.364 |
| 604 | 4 | 183 | 0 | 0 | 0 | 28.4 | 0.212 |
| 213 | 0 | 140 | 65 | 26 | 130 | 42.6 | 0.431 |
| 636 | 5 | 104 | 74 | 0 | 0 | 28.8 | 0.153 |
| 544 | 1 | 88 | 78 | 29 | 76 | 32.0 | 0.365 |
| 86 | 13 | 106 | 72 | 54 | 0 | 36.6 | 0.178 |
| 208 | 1 | 96 | 64 | 27 | 87 | 33.2 | 0.289 |
| 281 | 10 | 129 | 76 | 28 | 122 | 35.9 | 0.280 |
| 209 | 7 | 184 | 84 | 33 | 0 | 35.5 | 0.355 |
| 581 | 6 | 109 | 60 | 27 | 0 | 25.0 | 0.206 |
| 639 | 1 | 100 | 74 | 12 | 46 | 19.5 | 0.149 |
| 328 | 2 | 102 | 86 | 36 | 120 | 45.5 | 0.127 |
| 431 | 3 | 89 | 74 | 16 | 85 | 30.4 | 0.551 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 718 | 1 | 108 | 60 | 46 | 178 | 35.5 | 0.415 |
| 90 | 1 | 80 | 55 | 0 | 0 | 19.1 | 0.258 |
| 377 | 1 | 87 | 60 | 37 | 75 | 37.2 | 0.509 |

| Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction |
|-------------|---------|---------------|---------------|---------|-----|--------------------------|
| 235 | 4 | 171 | 72 | 0 | 0 | 43.6 |
| 158 | 2 | 88 | 74 | 19 | 53 | 29.0 |
| 69 | 4 | 146 | 85 | 27 | 100 | 28.9 |
| 260 | 3 | 191 | 68 | 15 | 130 | 30.9 |
| 131 | 9 | 122 | 56 | 0 | 0 | 33.3 |
| 44 | 7 | 159 | 64 | 0 | 0 | 27.4 |
| 70 | 2 | 100 | 66 | 20 | 90 | 32.9 |
| 264 | 4 | 123 | 62 | 0 | 0 | 32.0 |
| 673 | 3 | 123 | 100 | 35 | 240 | 57.3 |
| 286 | 5 | 155 | 84 | 44 | 545 | 38.7 |
| 640 | 0 | 102 | 86 | 17 | 105 | 29.3 |
| 135 | 2 | 125 | 60 | 20 | 140 | 33.8 |
| 745 | 12 | 100 | 84 | 33 | 105 | 30.0 |
| 165 | 6 | 104 | 74 | 18 | 156 | 29.9 |
| 164 | 0 | 131 | 88 | 0 | 0 | 31.6 |
| 28 | 13 | 145 | 82 | 19 | 110 | 22.2 |
| 608 | 0 | 152 | 82 | 39 | 272 | 41.5 |
| 583 | 8 | 100 | 76 | 0 | 0 | 38.7 |
| 746 | 1 | 147 | 94 | 41 | 0 | 49.3 |
| 292 | 2 | 128 | 78 | 37 | 182 | 43.3 |
| 136 | 0 | 100 | 70 | 26 | 50 | 30.8 |
| 432 | 1 | 80 | 74 | 11 | 60 | 30.0 |
| 554 | 1 | 84 | 64 | 23 | 115 | 36.9 |
| 319 | 6 | 194 | 78 | 0 | 0 | 23.5 |
| 594 | 6 | 123 | 72 | 45 | 230 | 33.6 |
| 6 | 3 | 78 | 50 | 32 | 88 | 31.0 |
| 615 | 3 | 106 | 72 | 0 | 0 | 25.8 |

192 rows × 8 columns

In [21]: `from sklearn import linear_model`

```
In [23]: logr = linear_model.LogisticRegression()
logr.fit(X,Y)
```

```
C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().  
y = column_or_1d(y, warn=True)
```

```
In [24]: y_pred=logr.predict(X_test)  
y pred
```

Step-3. [Predict on a new sample]

```
In [28]: new=logr.predict([[6,200,90,10,25,23.3,.672,42]])
if new==0:
    print("Non-diabetic patient",new)
else:
    print("Diabetic patient",new)
```

Diabetic patient [1]

Precision

```
In [31]: from sklearn.metrics import precision_score  
print(precision_score(y_test,y_pred))
```

9-7

Recall

```
In [32]: from sklearn.metrics import recall_score  
print(recall_score(y_test, y_pred))
```

0.6086956521739131

AUC Scores

```
In [33]: from sklearn.metrics import roc_auc_score
print(roc_auc_score(y_test,y_pred))
```

0.7311770943796395

```
In [37]: from sklearn import metrics
from sklearn.linear_model import LogisticRegression
```

```
In [39]: model = LogisticRegression(random_state=0, solver='lbfgs',multi_class='multinomial')
model.fit(X_train,y_train)
prediction=model.predict(X_test)
print('The accuracy of the Logistic Regression is', metrics.accuracy_score(prediction,y_test))
```

C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

The accuracy of the Logistic Regression is 0.7291666666666666

Step-4. [Understand Correlation]

```
In [34]: from sklearn.metrics import confusion_matrix
cfm=confusion_matrix(y_test,y_pred)
cfm
```

```
Out[34]: array([[105,  18],
       [ 27,  42]], dtype=int64)
```

```
In [36]: import seaborn as sns
sns.heatmap(cfm, annot=True)
```

```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x2d4d7f94cf8>
```



Step-5. [Normalization using MinMaxScaler and rebuild LOR]

```
In [40]: #Normalizing using MinMaxScaler
from sklearn.preprocessing import MinMaxScaler
mm=MinMaxScaler()
mm_X_train=mm.fit_transform(X_train)
mm_X_train
```

```
Out[40]: array([[0.76470588, 0.64824121, 0.           , ..., 0.59463487, 0.20964987,
   0.38333333],
   [0.23529412, 0.64824121, 0.70491803, ..., 0.52309985, 0.06532878,
   0.03333333],
   [0.17647059, 0.30653266, 0.67213115, ..., 0.51266766, 0.0704526 ,
   0.41666667],
   ...,
   [0.58823529, 0.50753769, 0.70491803, ..., 0.67958271, 0.45175064,
   0.28333333],
   [0.           , 0.70854271, 0.           , ..., 0.6318927 , 0.05422716,
   0.13333333],
   [0.           , 0.6281407 , 0.78688525, ..., 0.33532042, 0.07856533,
   0.        ]])
```

```
In [41]: mm_X_test=mm.transform(X_test)
mm_X_test
```

```
Out[41]: array([[0.35294118, 0.49246231, 0.47540984, ..., 0.50670641, 0.15029889,
   0.36666667],
   [0.11764706, 0.56281407, 0.6147541 , ..., 0.53204173, 0.02988898,
   0.        ],
   [0.11764706, 0.54271357, 0.52459016, ..., 0.45901639, 0.03415884,
   0.        ],
   ...,
   [0.35294118, 0.61809045, 0.59016393, ..., 0.50074516, 0.27967549,
   0.21666667],
   [0.17647059, 0.3919598 , 0.40983607, ..., 0.46199702, 0.07258753,
   0.08333333],
   [0.17647059, 0.53266332, 0.59016393, ..., 0.38450075, 0.05508113,
   0.1        ]])
```

```
In [42]: #ReBuild LOR Model
mm_lor=LogisticRegression()
mm_lor=mm_lor.fit(mm_X_train,y_train)
```

```
C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

```
In [43]: mm_y_pred=mm_lor.predict(mm_X_test)  
mm_y_pred
```

Accuracy

```
In [44]: model = LogisticRegression(random_state=0, solver='lbfgs',multi_class='multinomial')
model.fit(X_train,y_train)
prediction=model.predict(X_test)
print('The accuracy of the Logistic Regression is', metrics.accuracy_score(y_test,
```

The accuracy of the Logistic Regression is 0.7239583333333334

```
C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().  
    y = column_or_1d(y, warn=True)
```

Precision

```
In [45]: print(precision_score(y_test,mm_y_pred))
```

0.6538461538461539

Recall

```
In [46]: print(recall_score(y_test,mm_y_pred))
```

0.4927536231884058

AUC scores

```
In [47]: mm_auc=print(roc_auc_score(y_test,mm_y_pred))  
mm_auc
```

0.6732060798868859

Step-6: [Normalization using StandardScaler and rebuild LOR]

```
In [48]: from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
ss_X_train=ss.fit_transform(X_train)
ss_X_train
```

```
Out[48]: array([[ 2.80346794,  0.25977903, -3.78077929, ... ,  1.03974028,
   0.29608546,  0.96352088], [ 0.07832678,  0.25977903,  0.89724451, ... ,  0.40945373,
  -0.70087555, -0.86295593], [-0.22446668, -1.85825286,  0.67966201, ... ,  0.31753694,
  -0.66548048,  1.13747105], ... , [ 1.89508755, -0.61235174,  0.89724451, ... ,  1.78820556,
  1.96850229,  0.44167036], [-1.13284707,  0.63354937, -3.78077929, ... ,  1.36801453,
  -0.77756486, -0.34110542], [-1.13284707,  0.13518892,  1.44120077, ... , -1.24504846,
  -0.6094383 , -1.03690611]])
```

```
In [49]: ss_X_test=ss.transform(X_test)
ss_X_test
```

```
Out[49]: array([[ 0.6839137 , -0.70579433, -0.625833 , ... ,  0.26501306,
  -0.11390738,  0.87654579], [-0.52726014, -0.26972894,  0.29889263, ... ,  0.48823955,
  -0.94569142, -1.03690611], [-0.52726014, -0.39431905, -0.29945925, ... , -0.15517797,
  -0.91619553, -1.03690611], ... , [ 0.6839137 ,  0.07289387,  0.13570575, ... ,  0.21248918,
  0.77981801,  0.09377001], [-0.22446668, -1.32874488, -1.06099801, ... , -0.12891603,
  -0.65073254, -0.60203068], [-0.22446668, -0.45661411,  0.13570575, ... , -0.81172646,
  -0.77166568, -0.51505559]])
```

```
In [50]: #Rebuild LOR
ss_lor=LogisticRegression()
ss_lor.fit(ss_X_train,y_train)
ss_y_pred=ss_lor.predict(ss_X_test)
```

```
C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

Accuracy

```
In [51]: model = LogisticRegression(random_state=0, solver='lbfgs',multi_class='multinomial')
model.fit(X_train,y_train)
prediction=model.predict(X_test)
print('The accuracy of the Logistic Regression is', metrics.accuracy_score(y_test))
```

The accuracy of the Logistic Regression is 0.7291666666666666

```
C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

Precision

```
In [52]: print(precision_score(y_test,ss_y_pred))
```

0.6164383561643836

Recall

```
In [53]: print(recall_score(y_test,ss_y_pred))
```

0.6521739130434783

AUC scores

```
In [54]: auc_ss=print(roc_auc_score(y_test,ss_y_pred))
auc_ss
```

0.7122658183103571

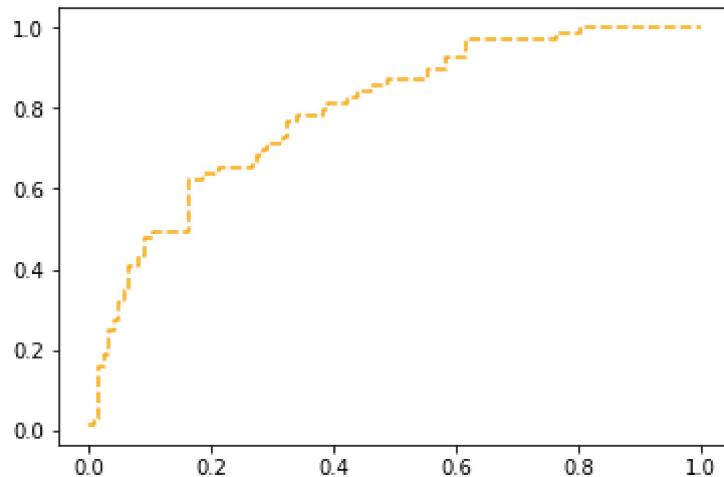
Step-7. [Plot ROC Curve]

```
In [55]: pred_prob1=mm_lor.predict_proba(mm_X_test)
```

```
In [56]: from sklearn.metrics import roc_curve
fpr1, tpr1, thresh1 = roc_curve(y_test, pred_prob1[:,1], pos_label=1)
```

```
In [57]: import matplotlib.pyplot as plt  
plt.plot(fpr1,tpr1,linestyle='--',color='orange',label='MinMaxScaler values')
```

```
Out[57]: [<matplotlib.lines.Line2D at 0x2d4d805d0f0>]
```



Step-8. [Comparison with KNN classifier]

```
In [58]: from sklearn.neighbors import KNeighborsClassifier  
knn=KNeighborsClassifier(n_neighbors=4)  
knn=knn.fit(X_train,y_train)
```

C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\ipykernel_launcher.py:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

This is separate from the ipykernel package so we can avoid doing imports until

```
In [59]: knn_y_pred=knn.predict(X_test)
```

```
In [60]: from sklearn.preprocessing import MinMaxScaler
m=MinMaxScaler()
m_X_train=m.fit_transform(X_train)
m_X_train
```

```
Out[60]: array([[0.76470588, 0.64824121, 0.38333333, 0.23529412, 0.64824121, 0.70491803, 0.52309985, 0.06532878, 0.17647059, 0.30653266, 0.67213115, 0.51266766, 0.0704526, 0.41666667, 0.58823529, 0.50753769, 0.70491803, 0.67958271, 0.45175064, 0.28333333], [0., 0.70854271, 0.13333333], [0., 0.6281407, 0.78688525, 0.33532042, 0.07856533, 0.]])
```

```
In [61]: m_X_test=m.transform(X_test)
m_X_test
```

```
Out[61]: array([[0.35294118, 0.49246231, 0.47540984, 0.50670641, 0.15029889, 0.36666667], [0.11764706, 0.56281407, 0.6147541, 0.53204173, 0.02988898, 0.], [0.11764706, 0.54271357, 0.52459016, 0.45901639, 0.03415884, 0.], 0.35294118, 0.61809045, 0.59016393, 0.50074516, 0.27967549, 0.21666667], [0.17647059, 0.3919598, 0.40983607, 0.46199702, 0.07258753, 0.08333333], [0.17647059, 0.53266332, 0.59016393, 0.38450075, 0.05508113, 0.1]])
```

```
In [62]: m_knn=KNeighborsClassifier()
m_knn=m_knn.fit(m_X_train,y_train)
```

```
C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\ipykernel_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
In [63]: m_y_pred=m_knn.predict(m_X_test)  
m_y_pred
```

Classification Metrics

Accuracy

```
In [64]: model = LogisticRegression(random_state=0, solver='lbfgs',multi_class='multinomial')
model.fit(X_train,y_train)
prediction=model.predict(X_test)
print('The accuracy of the Logistic Regression is', metrics.accuracy_score(y_test,
```

The accuracy of the Logistic Regression is 0.6770833333333334

C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

Precision

```
In [65]: print(precision_score(y_test, m_y_pred))
```

0.5522388059701493

Recall

```
In [66]: print(recall_score(y_test,m_y_pred))
```

0.5362318840579711

AUC Scores

```
In [67]: knn_auc=print(roc_auc_score(y_test,m_y_pred))  
knn_auc
```

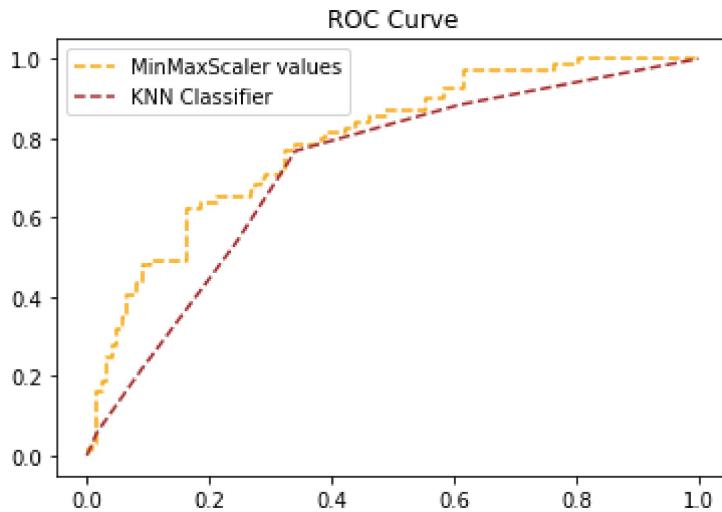
0.6461647225167905

Step-9. [Update ROC Curve]

```
In [68]: pred_prob2=m_knn.predict_proba(m_X_test)
```

```
In [69]: from sklearn.metrics import roc_curve
fpr2,tpr2,thresh2=roc_curve(y_test,pred_prob2[:,1],pos_label=1)
```

```
In [70]: import matplotlib.pyplot as plt
plt.plot(fpr1,tpr1,linestyle='--',color='orange',label='MinMaxScaler values')
plt.plot(fpr2,tpr2,linestyle='--',color='brown',label='KNN Classifier')
plt.title('ROC Curve')
plt.legend(loc='best')
plt.savefig('ROC',dpi=300)
plt.show()
```



Step-10. [Regularization]

```
In [71]: from sklearn.linear_model import LogisticRegressionCV
model1=LogisticRegressionCV(Cs=10,cv=4,penalty='l1',solver='liblinear')
model2=LogisticRegressionCV(Cs=10,cv=4,penalty='l2')
```

```
In [72]: model1.fit(mm_X_train,y_train)
model2.fit(mm_X_train,y_train)
```

```
C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

```
Out[72]: LogisticRegressionCV(Cs=10, class_weight=None, cv=4, dual=False,
fit_intercept=True, intercept_scaling=1.0, max_iter=100,
multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
refit=True, scoring=None, solver='lbfgs', tol=0.0001, verbose=0)
```

```
In [74]: rg_y_pred1 = model1.predict(mm_X_test)
rg_y_pred2 = model2.predict(mm_X_test)
```

AUC SCORE OF L1

```
In [75]: from sklearn.metrics import roc_auc_score
l1_auc = roc_auc_score(y_test, rg_y_pred1)
l1_auc = (' LOR L1 MINMAX AUC', l1_auc)
l1_auc
```

```
Out[75]: (' LOR L1 MINMAX AUC', 0.694591728525981)
```

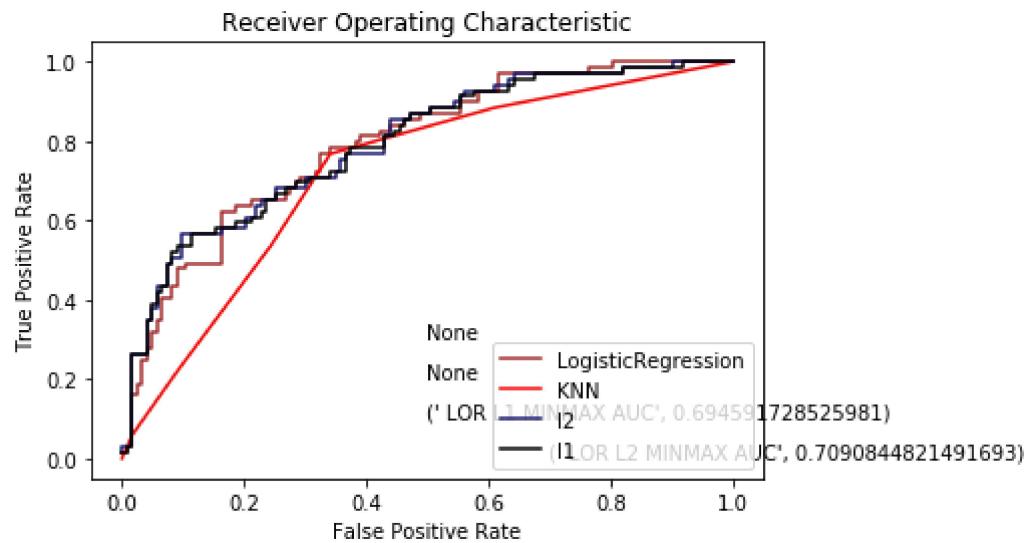
AUC SCORE OF L2

```
In [76]: from sklearn.metrics import roc_auc_score
l2_auc = roc_auc_score(y_test, rg_y_pred2)
l2_auc = (' LOR L2 MINMAX AUC', l2_auc)
l2_auc
```

```
Out[76]: (' LOR L2 MINMAX AUC', 0.7090844821491693)
```

STEP 12 : UPDATE ROC CURVE

```
In [77]: pred_prb7 = model1.predict_proba(mm_X_test)
pred_prb8 = model2.predict_proba(mm_X_test)
fpr,tbr,threshold = roc_curve(y_test, pred_prob1[:,1],pos_label=1)
fpr1,tbr1,threshold1 = roc_curve(y_test, pred_prob2[:,1],pos_label=1)
fpr2,tbr2,threshold2= roc_curve(y_test, pred_prb7[:,1],pos_label=1)
fpr3,tbr3,threshold3 = roc_curve(y_test, pred_prb8[:,1],pos_label=1)
plt.plot(fpr, tbr, linestyle='--', color='brown', label='LogisticRegression')
plt.plot(fpr1, tbr1, linestyle='--', color='red', label='KNN')
plt.plot(fpr3, tbr3, linestyle='--', color='midnightblue', label='l2')
plt.plot(fpr2, tbr2, linestyle='--', color='black', label='l1')
plt.annotate(xy=[0.5,0.3],s= auc_ss)
plt.annotate(xy=[0.5,0.2],s= knn_auc)
plt.annotate(xy=[0.5,0.1],s= l1_auc)
plt.annotate(xy=[0.7,0],s= l2_auc)
plt.title('Receiver Operating Characteristic')
plt.legend(loc = 'best')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [ ]:
```